

CV2 HW0

Qianhuang Li
ql2465

Code:

```
import cv2 as cv

# Store the height and width of the images
height, width = images[0].shape

# Pad each image with black by 30 pixels. You do not need to use this, but
# padding may make your implementation easier.

pad_size = 30
images_pad = [np.pad(x, pad_size, mode='constant') for x in images]

# Given two matrices, write a function that returns a number of how well they are aligned.
# The lower the number, the better they are aligned. There are a variety of scoring functions
# you can use. The simplest one is the Euclidean distance between the two matrices.

def score_function(im1, im2):

    score = np.linalg.norm(im1 - im2)

    return score

# Given two matrices chan1 and chan2, return a tuple of how to shift chan2 into chan1. This
# function should search over many different shifts, and find the best shift that minimizes
# the scoring function defined above.

def align_channels(chan1, chan2):

    best_offset = (0,0)

    best_score = 1000000

    index = 0

    for i in range(-30,30):
        for j in range(-30,30):
            index +=1

            offset = (i, j)

            chan1_n = chan1[200:600, 300:600]

            chan2_n = chan2[(200 - i):(600 - i), (300 - j):(600 - j)]

            score = score_function(chan1_n, chan2_n)

            if best_score > score:
```

```

        best_score = score

        best_offset = offset

    print(index, " ",best_score)

    return best_offset

rg_dy, rg_dx = align_channels(images_pad[0], images_pad[1])
rb_dy, rb_dx = align_channels(images_pad[0], images_pad[2])
print(rg_dy, rg_dx,rb_dy, rb_dx)

# Use the best alignments to now combine the three images. You should use any of the variables
# above to return a tensor that is (Height)x(Width)x3, which is a color image that you can visualize.

def combine_images():

    #TODO

    obj = np.float32([[1, 0, rg_dx], [0, 1, rg_dy]])
    image_g = cv.warpAffine(images[1], obj, (width, height))

    obj = np.float32([[1, 0, rb_dx], [0, 1, rb_dy]])
    image_b = cv.warpAffine(images[2], obj, (width, height))

    final_image = np.stack((images[0],image_g,image_b),axis = -1)

    return final_image

final_image = combine_images()

if final_image is not None:


    show_image(final_image)

```

Results:

I got two similar results:

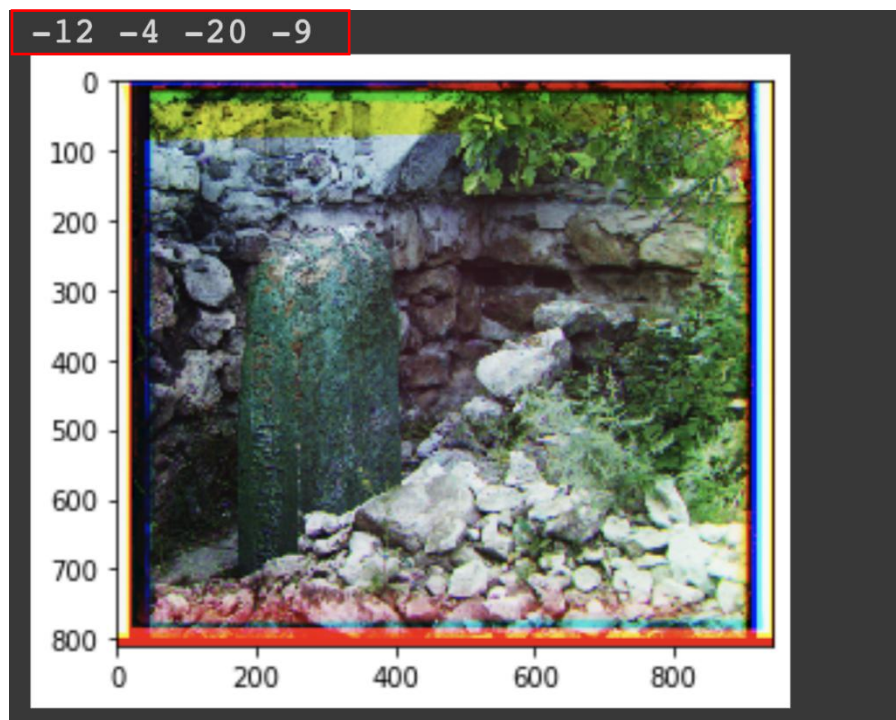
The first one used the central area from [200:600, 300:600] in Red Image, it took two seconds to complete the task.



Picture1 Upper line is the shift pixels of Green and Blue Images

Below is the combined image

The second one used the central area from [200:500, 300:500] in Red Image, it took one seconds to complete the task.



Picture1 Upper line is the shift pixels of Green and Blue Images

Below is the combined image

My colab link: <https://colab.research.google.com/drive/1e-8ByD2b-M1ZLGcKd-8KX3WZdxoWqlmm?usp=sharing>